

Design and Evaluation of an Edge Concurrency Control Protocol for Distributed Graph Databases

Paul Ezhilchelvan¹, Isi Mitrani¹, Jack Waudby¹ & Jim Webber²

November 28, 2019

¹Newcastle University

²Neo4j



Graph Databases

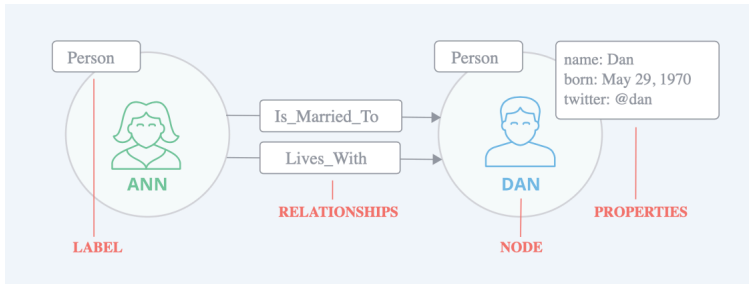


Figure 1: Labeled property graph¹

¹<https://neo4j.com>

- **Efficient graph analysis:** reachability, pattern matching, shortest path search and clustering

- **Efficient graph analysis:** reachability, pattern matching, shortest path search and clustering
- **Use cases:** telecommunications, pharma, publishing, finance, social media

Graph Database Popularity




Figure 2: Database popularity by data model²

¹https://db-engines.com/en/ranking_categories

- Graph exceeds the storage capacity of a single machine

- Graph exceeds the storage capacity of a single machine
- Partition the graph across multiple machines in a cluster, with partitions replicated for fault-tolerance and availability

- Graph exceeds the storage capacity of a single machine
- Partition the graph across multiple machines in a cluster, with partitions replicated for fault-tolerance and availability
- Graph partitioning inevitably leads to **distributed edges**




Figure 3: Connected graph

Distributed Graph Databases: Partitioning




Figure 4: Graph partitioned across 3 servers

A logical edge is represented by 2 physical records



Figure 5: A reciprocal consistent edge

Reciprocal Consistency

A logical edge is represented by 2 physical records




Figure 5: A reciprocal consistent edge

Reciprocal Consistency Violation

Two concurrent transactions:

- Mr Red requests to book the seat, T_R
- Mr Blue requests to book the seat, T_B

Reciprocal Consistency Violation




Figure 6: Concurrent transactions interleave and violate reciprocal consistency

Reciprocal Consistency Violation




Figure 6: Concurrent transactions interleave and violate reciprocal consistency

Reciprocal Consistency Violation




Figure 6: Concurrent transactions interleave and violate reciprocal consistency

Reciprocal Consistency Violation




Figure 6: Concurrent transactions interleave and violate reciprocal consistency

Reciprocal Consistency Violation




Figure 6: Concurrent transactions interleave and violate reciprocal consistency

Reciprocal Consistency Violation




Figure 7: Reciprocal inconsistent distributed edge

Reciprocal Consistency Violation




Figure 7: Reciprocal inconsistent distributed edge

Distributed Edge Reciprocal Consistency

- Without concurrency control, distributed edges can become reciprocally inconsistent
 - A distributed edge's reciprocal information is separated by the network

³Ezhilchelvan, P. et al, *On the degradation of distributed graph databases with eventual consistency*. European Performance Engineering Workshop 2018.

Distributed Edge Reciprocal Consistency

- Without concurrency control, distributed edges can become reciprocally inconsistent
 - A distributed edge's reciprocal information is separated by the network
- Reciprocal inconsistency is a source of corruption³

³Ezhilchelvan, P. et al, *On the degradation of distributed graph databases with eventual consistency*. European Performance Engineering Workshop 2018.

Distributed Edge Reciprocal Consistency

- Without concurrency control, distributed edges can become reciprocally inconsistent
 - A distributed edge's reciprocal information is separated by the network
- Reciprocal inconsistency is a source of corruption³
- No known concurrency control protocol exists specific to graph databases and maintaining reciprocal consistency

³Ezhilchelvan, P. et al, *On the degradation of distributed graph databases with eventual consistency*. European Performance Engineering Workshop 2018.

Our Solution: Collision Detection

- Carefully abort transaction(s)

Our Solution: Collision Detection

- Carefully abort transaction(s)
- No centralized control or synchronized clock

Our Solution: Collision Detection


- Carefully abort transaction(s)
- No centralized control or synchronized clock
- Permits interleavings that preserve reciprocal consistency

- Updates are initially provisional


- Updates are initially provisional
- Transactions distinguish between their first and second updates using labels

- Updates are initially provisional
- Transactions distinguish between their first and second updates using labels
- When a transaction attempts to update a given record, it can identify all other transactions that have earlier updated that record provisionally


Interference-free Update




Interference-free Update



Interference-free Update




Interference-free Update



Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort


Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort




Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort



Collision Detection

Rule: For any “1” seen, there must be a “2” in the opposite end. Else, abort



A second consistency problem related to distributed edges:

A second consistency problem related to distributed edges:

- Consider 2 transactions that update 2 or more of the *same* distributed edges

A second consistency problem related to distributed edges:

- Consider 2 transactions that update 2 or more of the *same* distributed edges
- Each update is reciprocally consistent

A second consistency problem related to distributed edges:

- Consider 2 transactions that update 2 or more of the *same* distributed edges
- Each update is reciprocally consistent
- Then if T_R updates before T_B on a given edge then this order should be preserved across all edges

Edge-Order Consistency Violation

Consider two edges,




Figure 8: Two distributed edges

Two concurrent transactions by a travel agent:

- Requests to book room and seat for Mr Red, T_R
- Requests to book room and seat for Mr Blue, T_B

Edge-Order Consistency Violation




Figure 9: Edge-order consistency violation

Edge-Order Consistency Violation




Figure 9: Edge-order consistency violation

Edge-Order Consistency Violation




Figure 9: Edge-order consistency violation

Edge-Order Consistency Violation




Figure 9: Edge-order consistency violation

Edge-Order Consistency Violation




Figure 9: Edge-order consistency violation

Edge-Order Consistency Violation




Figure 9: Edge-order consistency violation

Edge-Order Consistency Violation




Figure 10: Edge-order consistency violation

- Transactions collect predecessors from each update

Our Solution: Order Arbitration

- Transactions collect predecessors from each update
- Centralized arbiter maintains some global state


Our Solution: Order Arbitration

- Transactions collect predecessors from each update
- Centralized arbiter maintains some global state
- Global state used to detect edge-order violations

Our Solution: Order Arbitration


- Transactions collect predecessors from each update
- Centralized arbiter maintains some global state
- Global state used to detect edge-order violations
- Offending transactions are aborted

Rule: Collect predecessors after every successful update





Order Arbitration

Rule: Collect predecessors after every successful update





Order Arbitration

Rule: Collect predecessors after every successful update



Order Arbitration


Rule: Collect predecessors after every successful update




Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list

Arbiter




Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



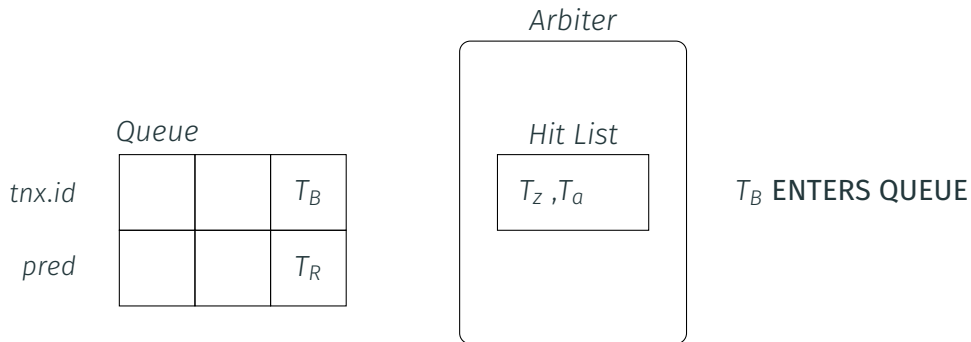
Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



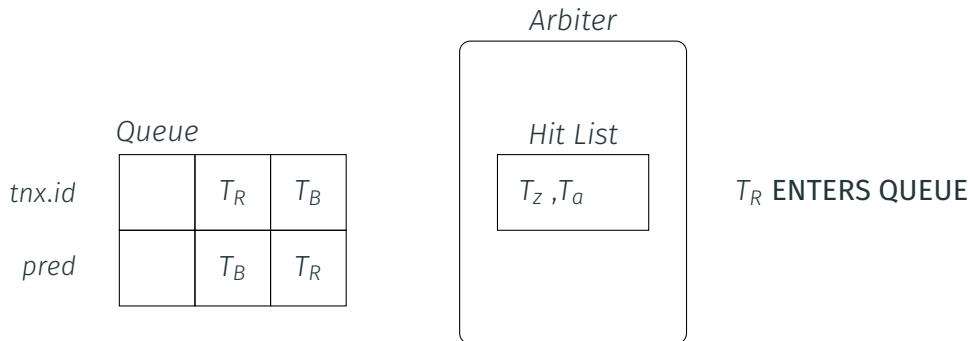
Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



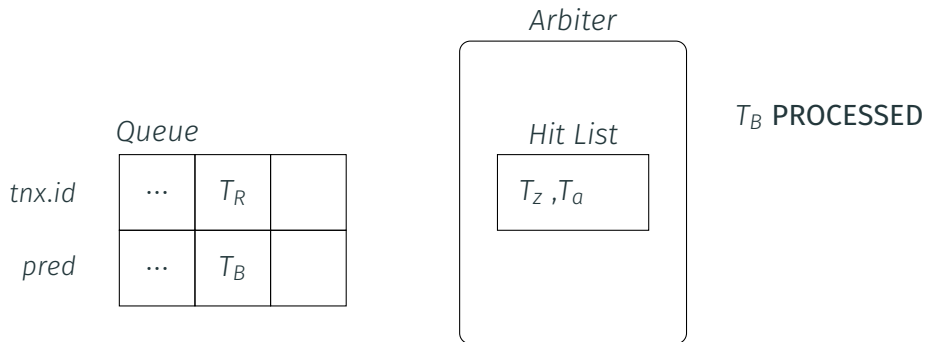
Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



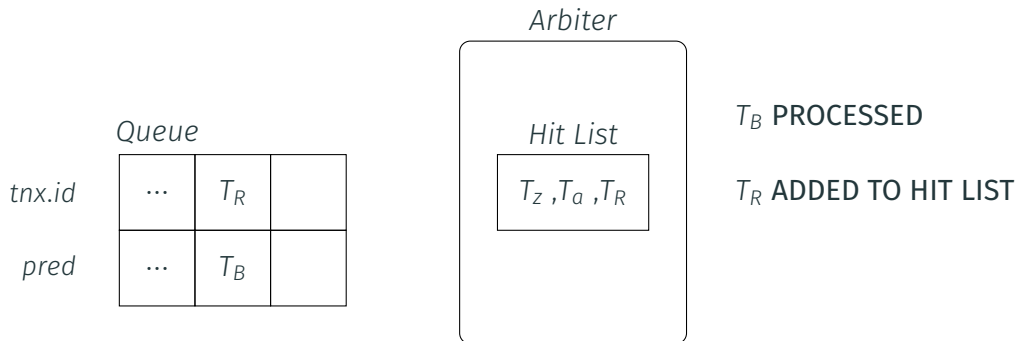
Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



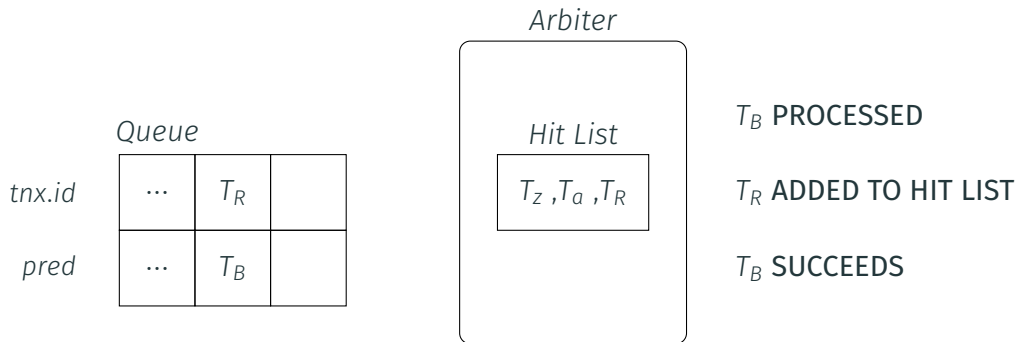
Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



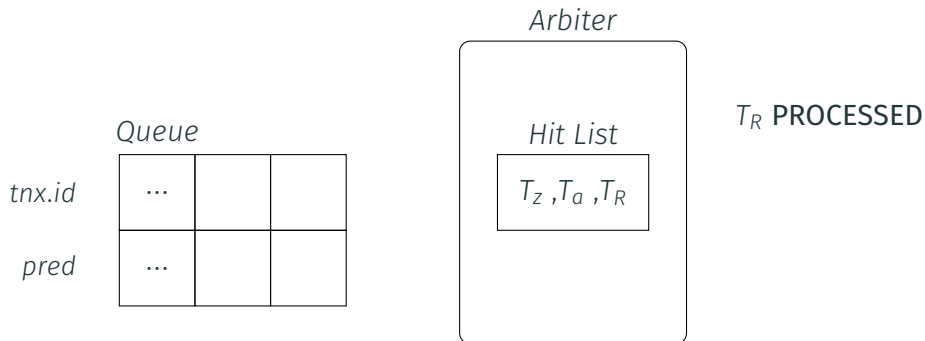
Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



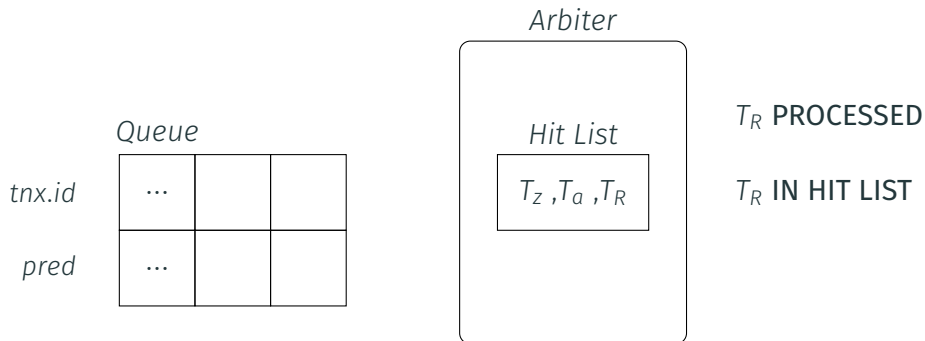
Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



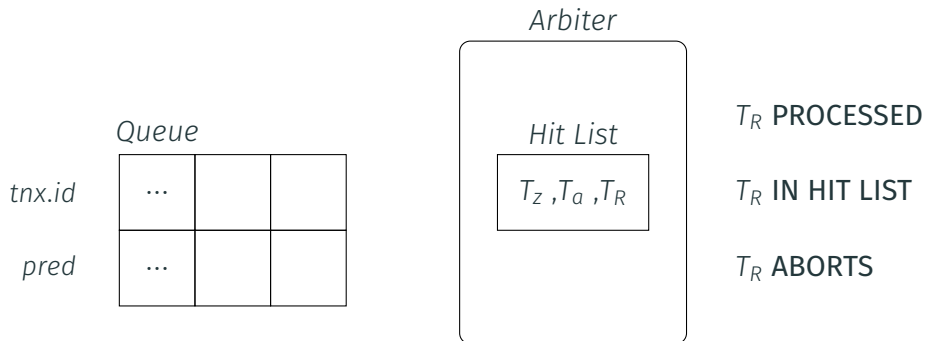
Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



Order Arbitration

Rule: If transaction exists in hit list, abort. Else, merge predecessors into hit list



If a transaction updates 1 or more distributed edges:

- Collect all predecessors

If a transaction updates 1 or more distributed edges:

- Collect all predecessors
- Go to arbiter

If a transaction updates 1 or more distributed edges:

- Collect all predecessors
- Go to arbiter
- If not in hit list then proceed and enter predecessors into the hit list

If a transaction updates 1 or more distributed edges:

- Collect all predecessors
- Go to arbiter
- If not in hit list then proceed and enter predecessors into the hit list
- Else in hit list then abort

1. Collision detection: enforces **reciprocal consistency** for distributed edges
2. Order arbitration: enforces **edge-order consistency** between transactions

Performance measures of interest:

- Average number of transactions that are aborted
- Load at the arbiter

1. Developed approximate model
2. Measure accuracy of model through simulation

- Database size
- Transaction arrival rate
- Average updates per transaction
- Average arbiter service time
- Average network delay

Aborts per second (R) vs Database Size (N)




Figure 11: TPS = 1000, av updates/tnx = 5, av arbiter service time = 10ms, av network delay = 5ms

Aborts per second (R) vs Transaction Arrival Rate (λ)

Arbiter queue unstable at ~ 1500 TPS




Figure 12: size = 10K, av updates/tnx = 5, av arbiter service time = 10ms, av network delay = 5ms

Aborts per second (R) vs Transaction Arrival Rate (λ)

Arbiter queue unstable at ~ 1100 TPS




Figure 13: size = 10K, av updates/tnx = 5, av arbiter service time = 10ms, av network delay = 10ms

Aborts per second (R) vs Transaction Arrival Rate (λ)

Arbiter queue unstable at ~ 550 TPS




Figure 14: size = 10K, av updates/tnx = 10, av arbiter service time = 10ms, av network delay = 5ms

- Model accuracy similar to simulations under a variety of parameter settings

- Model accuracy similar to simulations under a variety of parameter settings
- Between 1-4% transactions are aborted

- Model accuracy similar to simulations under a variety of parameter settings
- Between 1-4% transactions are aborted
- Improve accuracy of simulation

Thanks for listening!

Any Questions?

Email: j.waudby2@newcastle.ac.uk

Twitter: [@waldberry_7](https://twitter.com/waudberry_7)